

Architectural Blueprints and Engineering Requirements for Next-Generation Low-Code Workflow Orchestration Engines

Comparative Analysis of Low-Code Orchestration Frameworks

The modern enterprise ecosystem relies heavily on visual workflow orchestrators to coordinate API integrations, process automated data pipelines, and deploy generative artificial intelligence agents¹. Modern software system design has shifted from writing brittle, hardcoded integration scripts to using visual programming languages that model automated processes as directed graphs². At the core of this transformation are open-source and fair-code platforms that balance developer flexibility with intuitive, drag-and-drop interfaces². The primary systems driving this space are n8n, Activepieces, and Windmill, each approaching the orchestration challenge with distinct architectural assumptions, performance guarantees, and execution models⁵.

The fair-code platform n8n relies on a comprehensive, pre-built library of over 1,300 nodes⁶ to model integrations³. Its interface allows developers to customize inputs using inline JavaScript expressions or write dedicated scripts within a centralized Code node⁸. The platform balances accessibility with developer control, supporting advanced multi-step logical structures like conditional branching (IF and Switch nodes), multi-system coordination, and error recovery workflows³.

Activepieces is an MIT-licensed alternative that focuses on non-technical business teams and citizen developers⁵. By organizing integrations as decoupled npm modules known as "pieces"⁶, the platform simplifies community contributions and keeps individual workflow sizes compact⁶. Activepieces uses dynamic forms to auto-resolve target APIs and provides visual vertical layout tools that remove the manual alignment work often associated with open-ended canvases¹¹.

Windmill is a developer-focused platform that uses a code-first, long-running orchestration architecture⁵. Rather than forcing developers to build logical operations using visual nodes, Windmill allows scripts to be written in standard languages like Python, TypeScript, Go, Rust, or Bash⁷. The system parses the script's parameters and auto-generates a user interface, API endpoints, and cron triggers¹³. This allows organizations to build complex, data-intensive pipelines without proprietary software development kits (SDKs) or vendor lock-in¹³.

Beyond these three systems, other notable alternatives provide distinct trade-offs¹⁵. Node-RED remains highly popular for Internet of Things (IoT) integrations, offering over 5,700


community-contributed nodes⁶. Visual-first systems like Make and Zapier cater to rapid SaaS automation with large pre-built catalogs, but they run on closed-source, cloud-hosted platforms that limit on-premise deployments and deep data governance³. Specialized platforms like Vellum focus on prompt engineering, model comparison, and systematic testing, but they lack the generic system-integration features found in broader orchestration engines¹⁵.

Table 1: Core Platform Attributes and Capabilities

Platform	License Model	Primary Languages	Focus Area	Integration Library Size	Notable Architectural Feature
n8n	Sustainable Use & Enterprise ⁹	JavaScript, Python (via Pyodide) ⁸	Advanced visual-first AI & API workflows ³	1,300+ Nodes ⁶	Visual representation of LangChain graph structures ⁹ .
Activepieces	MIT License (Community Edition) ⁵	TypeScript, JavaScript ¹¹	Non-technical business operations ⁵	580+ Pieces ⁶	Decoupled pieces compiled as single KB-scale files served over S3 ¹² .
Windmill	AGPL v3 (Community Edition) ⁵	Python, TypeScript, Go, Rust, Bash, PHP, SQL ⁷	Developer scripts, data pipelines & internal tools ⁷	Custom script scripts & community hub integrations ⁷	In-memory execution suspend and resume via state checkpoints ²⁰ .
Node-RED	Apache 2.0 ⁶	JavaScript ⁶	IoT telemetry & physical hardware automation ⁶	5,700+ Nodes ⁶	Browser-based visual wiring of JSON streams ⁶ .

Monorepo Architecture and Software Components

Building a scalable workflow engine requires a structured, multi-tier codebase that cleanly separates the visual interface, backend API routing, job queue dispatching, and execution environment²². A monorepo structure utilizing Turborepo provides a clean layout for managing these dependencies²³. The following architecture outlines the monorepo structure used by Activepieces, which serves as a model for developer contributions and code isolation²³.



At the user interface layer, performance depends on how efficiently the DOM handles thousands of nodes and connections²⁴. The visual frontend is typically built using React, using state-management libraries like React Query to handle server state²³. By grouping query hooks (such as `useQuery` and `useMutation`) inside a centralized module like `features/lib/feature-hooks.ts`, developers can decouple API communications from component UI code²⁵. Centralizing query keys and managing dialog states through a global store prevents prop drilling and ensures consistent UI rendering across the application canvas²⁵. For database communications, using an Object-Relational Mapper (TypeORM) simplifies migrations across database engines²². While SQLite or embedded databases like PGLite work well for local development, production environments require a highly concurrent relational database like PostgreSQL²².

In multi-tenant cloud platforms, the database schema must enforce absolute logical isolation¹⁸. The server layer must use safe factory structures to inject tenant contexts, ensuring that every SQL query dynamically filters data using the target `projectId` or `platformId`²⁸.

Database Schema and tenanted Data Tables

To maintain operational integrity and state tracking, the core database schema must support

tenant partitioning, workflow history versioning, credentials management, and detailed execution logs²². Additionally, enterprise platforms often require local datastores—such as n8n's native Data Tables—to let workflows store and reference structured data without setting up external databases²⁹.

In n8n, Data Tables are visually managed within the UI, are accessible programmatically via the /datatables REST API, and are capped at a default limit of 50MB per instance to prevent database bloat²⁹. The platform triggers warning flags when a tenant reaches 80% of this limit²⁹. Direct programmatic access from Code nodes is blocked to protect the primary database thread, requiring workflows to interact with tables using dedicated Data Table nodes²⁹. Developers often use schema-migration patterns (such as dynamic "create column if not exists" execution logic) to let AI agents or workflows dynamically update database schemas when handling incoming payloads with varying structures³⁰.

Table 2: Relational Database Schema Design

Table Name	Column Name	Data Type	Key Constraints	Description & Indexes
project [cite: 26]	id	UUID	PRIMARY KEY	Unique project workspace ID ²⁶ .
	name	VARCHAR(255)	NOT NULL	Workspace display name ²⁶ .
	created	TIMESTAMP	DEFAULT NOW()	Record creation date.
user [cite: 26]	id	UUID	PRIMARY KEY	Unique user ID ²⁶ .
	email	VARCHAR(255)	UNIQUE, NOT NULL	Authentication identifier ³¹ .
	password_hash	VARCHAR(255)	NOT NULL	Securely hashed password.
workflow_entity [cite: 26]	id	UUID	PRIMARY KEY	Unique workflow ID ³² .

	project_id	UUID	FOREIGN KEY	References project.id (Indexed for fast lookups) ²⁶ .
	name	VARCHAR(255)	NOT NULL	Name of the automation workflow ³² .
	active	BOOLEAN	DEFAULT FALSE	Execution status of triggers ³³ .
	nodes	JSONB	NOT NULL	Array of node structures, credentials, and parameters ³² .
	connections	JSONB	NOT NULL	Edge connectivity mappings ³² .
workflow_history [cite: 26]	id	UUID	PRIMARY KEY	Unique version identifier ²⁶ .
	workflow_id	UUID	FOREIGN KEY	References workflow_entity.id ²⁶ .
	nodes	JSONB	NOT NULL	Workflow node configuration snapshot ²⁶ .
	connections	JSONB	NOT NULL	Workflow edge layout snapshot.
	created_at	TIMESTAMP	DEFAULT NOW()	Timestamp of the change.

credentials_entity [cite: 26]	id	UUID	PRIMARY KEY	Unique credential ID ³² .
	name	VARCHAR(255)	NOT NULL	Name of the credential ³² .
	type	VARCHAR(128)	NOT NULL	Connector type identifier ³³ .
	data_encrypted	TEXT	NOT NULL	Encrypted credential data ²² .
	project_id	UUID	FOREIGN KEY	References project.id ²⁶ .
execution_entity [cite: 26]	id	UUID	PRIMARY KEY	Unique execution run ID ²² .
	workflow_id	UUID	FOREIGN KEY	References workflow_entity.id (Indexed) ⁹ .
	status	VARCHAR(64)	NOT NULL	Execution status (e.g., SUCCESS, FAILED) ²² .
	started_at	TIMESTAMP	NOT NULL	Timestamp when execution began.
	finished_at	TIMESTAMP	NULL	Timestamp when execution completed.

execution_data [cite: 26]	id	UUID	PRIMARY KEY	Unique execution data payload ID ²⁶ .
	execution_id	UUID	FOREIGN KEY	References execution_entity.id ²⁶ .
	step_data_json	JSONB	NOT NULL	Step-by-step inputs, outputs, and variables ²² .
webhook_entity [cite: 26]	id	UUID	PRIMARY KEY	Unique webhook identifier ²⁶ .
	workflow_id	UUID	FOREIGN KEY	References workflow_entity.id ²⁶ .
	webhook_path	VARCHAR(255)	UNIQUE	Unique path routing incoming requests ¹² .

Distributed Process Orchestration and Queue Mode Engineering

When running a workflow engine in a single-process container, the Node.js main thread handles backend APIs, incoming webhooks, cron scheduling, and CPU-heavy tasks like large data transformations simultaneously³⁵. Under heavy loads, such as a sudden surge of hundreds of concurrent webhook requests, this architecture can hit a single-process bottle-neck³⁶. The event loop can lag, causing webhook timeouts, slow editor loading times, and delayed executions¹.

To scale beyond these single-process limits, production environments use Queue Mode³⁵. This separates n8n or Activepieces into independent, decoupled services coordinated by a high-throughput Redis database running BullMQ¹.



The services are split into distinct operational roles³⁵:

- **Main Process:** Exposes ports to host the user interface, serves REST API calls, tracks schedules, and manages the WebSocket gateway³⁵. When a workflow runs, the main process creates an execution ID in PostgreSQL but delegates the actual execution by pushing a job to Redis²².
- **Webhook Processors:** Lightweight, dedicated processes optimized solely to receive incoming webhooks³⁶. They bypass workflow execution entirely, immediately writing payloads to the Redis queue to ensure fast response times¹.
- **Redis Queue Engine (BullMQ):** A high-speed memory layer that tracks active job queues, handles worker timeouts, and manages state locks³⁶.
- **Worker Fleet:** Stateless containers that pull execution IDs from Redis³⁵. The workers load the corresponding workflow structure and credentials from PostgreSQL, run the execution steps, and write the final results back to PostgreSQL³⁵.

Setting up a distributed fleet requires careful environment configuration³⁶. The main process and workers must use an identical encryption key (N8N_ENCRYPTION_KEY)³⁶. If keys mismatch, workers cannot decrypt credential packages, causing silent workflow failures³⁶.

Memory bloat is a common issue with BullMQ because completed job metadata is retained in Redis by default³⁶. To prevent high memory usage, developers configure the QUEUE_BULL_JOB_OPTIONS_REMOVE_ON_COMPLETE=1000 environment variable to prune historical metadata³⁶.

Additionally, setting N8N_GRACEFUL_SHUTDOWN_TIMEOUT=30 ensures workers complete active executions before shutting down, allowing zero-downtime rolling deployments³⁶.

Table 3: BullMQ Job Lifecycle States and Transitions

Job State	Target Service	Redis Mechanism	Transition Cause	Resource Impact
WAITING [cite: 35]	Redis / Ingestion Layer	Redis List / Hash ³⁷	Incoming webhook pushes job payload; waiting for worker polling ³⁵ .	Minimal: Memory consumption scales with queue depth ³⁵ .
ACTIVE [cite: 35]	Worker Fleet ³⁵	Redis Sorted Set with locks ³⁷	Worker claims job and starts processing workflow logic ³⁵ .	CPU usage spikes based on the complexity of node executions ³⁶ .

SUSPENDED [cite: 20]	Database / Redis	Persistent Storage Snapshot ²⁰	Workflow hits a sleep delay or waits for manual human approval ¹² .	Low: Worker threads are released; memory state is written to disk ²⁰ .
COMPLETED	Redis / Database	Pruned according to retention rules ²⁷	Workflow finishes successfully; execution records are committed ³⁵ .	Memory is freed; results are saved to PostgreSQL ³⁵ .
FAILED	Redis / Database	Dead-letter list queue ³	Execution encounters an error, exhausts retry attempts, or worker crashes ³ .	High database writes for logging error states and stack traces ²² .

This stateless worker architecture introduces autoscaling challenges on serverless runtimes like Google Cloud Run³⁹. Standard serverless autoscalers monitor incoming HTTP traffic to scale instances up or down³⁹. Because background workers pull jobs from Redis instead of receiving direct HTTP requests, standard autoscalers are "blind" to queue depth³⁹. This can cause worker instances to scale down to zero even while hundreds of heavy jobs are queuing in Redis³⁹. Developers can address this by configuring workers with a strict resources limit (e.g., 1 vCPU and 1GB RAM) and running a continuous lightweight health-ping container¹². When heavy workflow runs load the worker, the resource-constrained container quickly hits 100% CPU usage³⁹. This triggers the host system's CPU-based autoscaling rules, spinning up more worker instances to handle the load³⁹.

Security Engineering and Sandboxing Execution Jails

To prevent arbitrary script execution from turning the workflow engine into an open Remote Code Execution vector, system architects must implement strict, multi-layered isolation barriers⁴⁰.

Sandbox Escape Vector Analysis

The JavaScript runtime environment in Node.js is flexible, but its prototype-based inheritance model makes in-process sandboxing challenging to secure⁴⁰. Prototype pollution occurs when an execution path modifies properties on shared parent prototypes like `Object.prototype` or `Function.prototype`⁴⁰. This modification propagates across all objects in the runtime, allowing a

sandboxed script to escape its context and access host-level functions⁴¹.

Vulnerabilities like CVE-2026-25049 demonstrate how attackers can bypass static Abstract Syntax Tree (AST) analyzers⁴⁰. AST parsers scan input code for forbidden properties like `constructor` or `__proto__` to block malicious actions⁴⁰.

However, attackers can hide their intent using calculated properties and proxy objects⁴⁰. The AST parser flags the code as safe, but once evaluated, a proxy getter is triggered to escape the context⁴⁰:

```
JavaScript
// Demonstration of hiding malicious intent from AST Sanitizers
const calculatedExploit =
  "get" + "ter"
  "con" + "structor"
  "pro" + "totype"
  pollutedProperty true

```

Other attack vectors target execution contexts directly¹⁷. CVE-2025-68613 demonstrates how expression injection can expose the process object during inline evaluation, allowing direct access to the host's terminal tools¹⁷:

```
JavaScript
// Escape expression using the host process constructor
(function()
  var hostProcess = this.constructor.constructor("return process")
  return hostProcess.mainModule.require("child_process").execSync("id").toString()

```

Additionally, CVE-2026-21858 exploits Content-Type confusion during HTTP requests, allowing

attackers to perform arbitrary file reads¹⁷. This allows them to retrieve local credentials and database keys to forge administrator-level JSON Web Tokens (JWTs)¹⁷.

CVE-2026-22709 targets vm2 sandboxes by exploiting inconsistencies in how promise prototypes are handled in local and host contexts, bypassing promise callback sanitization entirely⁴⁴.

Furthermore, shared-execution environments must manage memory exposure risks⁴⁶. A race condition in the Node.js buffer allocation logic can disable the zero-fill toggle if execution is interrupted by a sandbox timeout⁴⁶. This can expose uninitialized heap memory, potentially revealing sensitive credentials, passwords, or API keys from previous execution threads⁴⁶.

Isolation Implementations

Modern platforms avoid in-process sandboxes and instead use isolated execution environments³⁸.



Architects use the following environments to enforce isolation:

- **V8 Isolates (via isolated-vm):** This library runs untrusted scripts in separate V8 instances, each with its own heap and thread pools¹⁸. It prevents scripts from accessing host-level constructors or properties⁴³. However, passing callbacks and tool states across the isolate boundary is complex and can introduce latency⁴⁸.
- **Independent Sandboxed Processes (SANDBOX_PROCESS):** Runs scripts in isolated processes managed by Linux namespaces, chroot directories, and memory-constrained

control groups³⁸.

- **WebAssembly WASI Environments:** Compiles custom actions into WebAssembly binaries (.wasm)¹⁸. The WasmEdge runtime enforces strict memory limits and blocks access to the file system and network unless explicitly allowed¹⁸.


In Activepieces, developers configure these sandbox behaviors using dedicated environment variables:

- `AP_EXECUTION_MODE`: Selects the sandboxing strategy (UNSANDBOXED, SANDBOX_PROCESS, SANDBOX_CODE_ONLY, or SANDBOX_CODE_AND_PROCESS)²⁷.
- `AP_SANDBOX_MEMORY_LIMIT`: Sets the maximum memory limit (in kilobytes) allocated to a sandboxed engine process (e.g., 1048576 for a 1GB limit)²⁷.
- `AP_FLOW_TIMEOUT_SECONDS`: Terminates runaway loops or long-running scripts by enforcing a maximum execution time limit (e.g., 600 seconds)²⁷.
- `AP_WORKER_CONCURRENCY`: Sets the number of concurrent executions processed by a single worker container, running each execution in its own sandbox instance²⁷.

Server-Side Request Forgery (SSRF) Mitigations

Because workflow nodes make automated HTTP calls using dynamic parameters, attackers can craft requests targeting internal IP addresses, local services, or cloud metadata endpoints (like AWS/GCP's 169.254.169.254)⁴⁰. This is known as a Server-Side Request Forgery (SSRF) attack⁴¹.





To secure outbound network traffic, platforms implement application-level network policies¹²:

- **n8n SSRF Egress Guard**: Setting `N8N_SSRF_PROTECTION_ENABLED=true` intercepts all requests before DNS resolution⁵⁰. The target hostname is checked against blocked ranges⁵⁰. If allowed, the engine performs a DNS lookup and validates the resolved IP address against configured blocklists (such as RFC 1918 private CIDR blocks or cloud metadata ranges) before opening an HTTP connection⁵⁰. Hostnames configured in `N8N_SSRF_ALLOWED_HOSTNAMES` bypass these IP blocklist checks to let workflows reach legitimate internal services⁵⁰.
- **Activepieces Network posturing**: Setting `AP_NETWORK_MODE=STRICT` installs an in-process interceptor that blocks outbound requests to localhost, loopback networks, and cloud metadata endpoints across all Node.js networking libraries (such as axios, fetch, undici, or raw TCP sockets)¹². Legitimate internal targets are allowed by adding their subnets (e.g., 10.10.0.0/24) to the `AP_SSRF_ALLOW_LIST`¹².

Algorithmic Logic of Graph Execution, Suspended States, and Waitpoints

At its core, a workflow engine is a graph execution system¹. It parses a workflow's node topology, validates its structure, and manages how data flows between execution steps¹.

Mathematical Graph Logic and Data Flows

Mathematically, a workflow is modeled as a Directed Acyclic Graph (DAG), represented as:

$$G = (V, E)$$

where V is the set of execution nodes and E is the set of directed edges routing data between nodes¹. Before running a workflow, the engine performs a topological sort to validate that the graph is acyclic¹:

$$\text{Topological Sort}(G) \implies \forall (u, v) \in E, \quad \text{index}(u) < \text{index}(v)$$

If the topological sort encounters a cycle, the validation fails and execution is aborted⁵³. Data flows between nodes as normalized JSON arrays²². This standard format ensures

consistency when passing data between different third-party APIs²²:

$$I_v = \mathcal{M}(O_u \mid (u, v) \in E)$$

where the input payload I_v of a node v is mapped ($\mathcal{M}$) from the outputs O_u of its parent nodes u ²².



Workflows support three primary execution patterns¹:

- **Sequential Flow:** Nodes run one after another in a linear sequence¹⁹.
- **Parallel Fan-Out / Fan-In:** The engine runs multiple independent branches concurrently and joins their outputs at a merge node before continuing execution¹⁹.
- **Sub-Workflows:** A parent workflow calls a child workflow to perform a sub-task¹. When calling sub-workflows, incoming data payloads are typically wrapped in an execution input object⁵⁵. This requires updating expressions within the sub-workflow to reference the new parent namespace (such as using `{{ $json.body.field }}` instead of `{{ $json.field }}`) to avoid mapping errors⁵⁵.

Durable Execution and Waitpoints

When a workflow runs, it often encounters steps that pause execution for extended periods,

such as waiting for a 2-day delay or a manual human approval gate¹¹. Standard execution engines that block threads during these delays waste system resources and risk data loss if the worker container restarts or crashes²⁰.

To resolve this issue, engines like Activepieces and Windmill use durable suspension models²⁰:

- **Activepieces Waitpoints:** When a workflow hits a waitpoint, the active process is halted¹². The current execution state, context variables, and completed outputs are written to the database²⁰. The job is removed from the active Redis queue, freeing up the worker's processing slot¹².
 - *Webhook Waitpoints:* Pause execution until a unique callback URL is called¹². When triggered, the callback's HTTP request body, query parameters, and headers are injected into the saved context¹². The engine then schedules the job on an available worker to resume execution¹².
 - *Delay Waitpoints:* Calculate a target resume timestamp¹². A scheduler monitors these timestamps and automatically enqueues the workflow for execution once the delay expires¹².
- **Windmill Workflows-As-Code:** Windmill allows developers to write workflow logic directly in TypeScript or Python²⁰. Its engine coordinates this execution using standard language constructs like loops and conditional functions²⁰:
 - Code logic is wrapped in a workflow() helper, and steps are annotated with the @task decorator²⁰.
 - The script executes step-by-step until it hits an uncached task(), sleep(), or waitForApproval() call²⁰.
 - The engine saves a checkpoint of all completed step results to the database and halts execution, releasing the worker²⁰.
 - When resuming, the script re-runs from the beginning²⁰. However, previously completed steps retrieve their outputs directly from the saved database checkpoints instead of executing again, allowing the script to continue from where it left off²⁰.
 - For lightweight expressions, Windmill uses the QuickJS engine, which provides 8-16x faster startup times than standard runtimes by lazily resolving resource variables on demand¹⁹.

Table 4: State Capture and Waitpoint Mechanisms

Feature	n8n	Activepieces	Windmill
Durable Resumption Method	Database-persisted execution states ²² .	Redis-backed BullMQ job queue suspension ³⁸ .	Database checkpoints combined with script re-runs ²⁰ .

Waitpoint Types	Webhooks, wait delays, and manual approval gates ⁸ .	Webhook callbacks and scheduled delay waitpoints ¹² .	sleep() delays and waitForApproval() human gates ²⁰ .
Worker Release Strategy	Releases execution workers when in queue mode ³⁵ .	Worker concurrency is set to 1, releasing slots during delays ¹² .	Script exits, fully releasing the worker container ¹³ .
Expression Evaluation	V8 VM context evaluation ⁴⁰ .	In-process TypeScript runtime ²³ .	QuickJS engine for fast expression evaluation ¹⁹ .
State Sharing	Global execution variables and context maps ⁸ .	Project-scoped variables and global context ²⁷ .	Custom flow states (setFlowUserState) ¹⁹ .

AI-First Architecture Patterns and Multi-Agent Orchestration

As visual orchestrators evolve, they are increasingly used as runtime engines for generative AI pipelines¹. Rather than simply calling an LLM endpoint, modern platforms integrate AI models with structured tools, system memory, and automated workflows¹⁶.





LangGraph Agent Orchestration and Streaming-First Principles

For real-time user experiences, AI workflow builders must stream responses rather than waiting for the complete LLM output³⁴. n8n uses a streaming-first architecture that streams LLM responses over HTTP using JSON-lines separated by custom Unicode markers³⁴. The backend coordinates complex task building using a multi-agent orchestration graph powered by LangGraph³⁴. The system divides task processing among six specialized agents³⁴:

- **Supervisor:** Acts as the central router, managing the flow of data and coordinating

sub-agents³⁴.

- **Discovery:** Identifies relevant node types and tools from the engine's registry³⁴.
- **Builder:** Configures node parameters and builds the dynamic JSON execution schemas³⁴.
- **Planner:** Drafts the execution steps and logical paths³⁴.
- **Responder:** Generates natural language summaries and structures output data³⁴.
- **Parameter Updater:** Handles updates to variables and tracks system state across runs³⁴.

These agents also integrate with Human-in-the-Loop (HITL) approval gates³⁴. If a planner agent suggests a sensitive action—such as executing a database migration or sending an email campaign—the system pauses execution³⁴. The workflow suspends and waits for administrator approval before running the suggested steps³⁴.

Table 5: 22-Step Execution Flow for Workflow Generation

Phase	Step	Action Event / Operation	Component Involved	Output Protocol / Persistence
Phase 1: User Input	1	User submits prompt in visual builder ³⁴ .	Frontend Canvas UI ³⁴	Emits @message event ³⁴ .
	2	Client triggers concurrency lock guards ³⁴ .	Frontend sendChatMessage() [cite: 34]	Prevents concurrent stream requests ³⁴ .
	3	Captures active canvas snapshot ³⁴ .	Canvas State Manager ³⁴	Used to revert changes if execution fails ³⁴ .
	4	Builds the context payload ³⁴ .	Client Engine ³⁴	Passes active node schemas to context ³⁴ .
	5	Dispatches HTTP POST request ³⁴ .	Client API Client ³⁴	Routes request to /ai/build endpoint ³⁴ .
Phase 2:	6	Authenticates	AiController.bu	Rejects

Backend		tenant and verifies licenses ³⁴ .	ild() [cite: 34]	request if limit is exceeded ³⁴ .
	7	Initializes builder service ³⁴ .	AiWorkflowBuilderService [cite: 34]	Resolves active model checkpoints ³⁴ .
	8	Instantiates LLM configuration ³⁴ .	Backend LLM Client ³⁴	Prepares security and API tokens ³⁴ .
	9	Evaluates target session checkpoint ³⁴ .	Session Controller ³⁴	Matches history ID mapping ³⁴ .
	10	Delegates routing to supervisor ³⁴ .	Supervisor Agent ³⁴	Routes execution to target subgraph ³⁴ .
	11	Scans node database for dependencies ³⁴ .	Discovery Agent ³⁴	Matches parameters against node types ³⁴ .
	12	Generates required node structures ³⁴ .	Builder Agent ³⁴	Builds step configuration fields ³⁴ .
	13	Translates node logic to JSON schema ³⁴ .	Planner Agent ³⁴	Formats output to fit DAG connections ³⁴ .
	14	Evaluates security parameters ³⁴ .	Egress Guardian ³⁴	Verifies SSRF security rules ⁵⁰ .
	15	Appends response text	Responder Agent ³⁴	Streams data using

		to output ³⁴ .		JSON-lines ³⁴ .
Phase 3: Response	16	Receives HTTP stream chunks ³⁴ .	Frontend Client ³⁴	Triggers client-side listener ³⁴ .
	17	Parses incoming stream ³⁴ .	onMessageUpdated() [cite: 34]	Converts JSON-lines to UI message arrays ³⁴ .
	18	Renders streamed text in chat window ³⁴ .	Chat UI Interface ³⁴	Updates assistant response in real time ³⁴ .
	19	Renders generated node layouts ³⁴ .	Canvas State Manager ³⁴	Dynamically adds nodes to visual editor ³⁴ .
	20	Signals completion of LLM stream ³⁴ .	stopStreaming() [cite: 34]	Closes active HTTP connection ³⁴ .
Phase 4: Persist	21	Starts a transaction to save the workflow ³⁴ .	WorkflowCreationService [cite: 34]	Validates output schema definitions ³⁴ .
	22	Commits workflow structures to DB ³⁴ .	Database Transaction layer ³⁴	Saves workflow versions to history tables ³⁴ .

Local Models and Model Context Protocol (MCP) Integration

While cloud-hosted models are convenient, privacy-conscious organizations often run open-source models locally using Ollama⁵⁷. This is configured using an HTTP Request node to route queries to local endpoints, keeping sensitive business data within the private network⁵⁷. Furthermore, modern workflow engines can expose their actions to external AI systems using Model Context Protocol (MCP) servers⁸. This allows external LLMs to interact with the engine programmatically, letting them trigger workflows, query database tables, and run automated

tasks directly from standard chat interfaces⁸.

Architectural Recommendations for Engineering a Next-Generation Orchestrator

To build a secure, highly scalable, and enterprise-grade workflow orchestration platform, developers should adhere to the following key design guidelines:

Implement an Asynchronous, Queue-Based Service Architecture

Decouple frontend servers and webhook receivers from the workflow execution layer³⁵. Use a lightweight framework like Fastify to handle incoming webhook requests and immediately write payloads to a Redis queue managed by BullMQ²². Run execution workloads on a stateless worker fleet that scales horizontally based on queue depth³⁵. This prevents slow, long-running workflows from blocking webhook ingestion or causing API lag³⁶.

Enforce Strict Sandbox Isolation for Untrusted Code

Avoid running user-submitted code in the primary host process, as in-process sandboxing is highly vulnerable to prototype pollution and sandbox escapes⁴⁰. Run JavaScript and Python scripts in isolated V8 isolates using isolated-vm or in separate processes secured by Linux namespaces and resource limits²⁷. Additionally, secure all outbound HTTP requests using a strict network proxy that intercept requests before DNS resolution to block access to private subnets and cloud metadata endpoints⁵⁰.

Use Durable State-Tracking for Long-Running Workflows

Do not keep worker threads open when workflows encounter delay steps or wait for manual human approvals²⁰. When hitting a delay, serialize the workflow's state, context variables, and execution history, write the snapshot to PostgreSQL under a SUSPENDED status, and release the worker back to the pool¹². Once the delay expires or an approval callback is received, reload the state snapshot and queue the workflow to resume execution on an available worker¹².

Structure the Codebase as a Modular, Tenanted Monorepo

Organize the project as a Turborepo-managed monorepo to cleanly separate shared models, custom pieces, visual builders, and worker engines²³. Use database frameworks like TypeORM to easily manage schema migrations across relational databases like PostgreSQL²². Ensure all database calls dynamically append tenant scoping variables (such as filtering by projectId or platformId) to prevent data leakage in multi-tenant deployments²⁸.

Build a Structured Data Exchange Layer for AI Workflows

Optimize the visual frontend to support real-time data streaming by formatting LLM responses as HTTP JSON-lines³⁴. Build structured interfaces that allow multi-agent coordination using LangGraph routing patterns, and isolate AI actions from the underlying integrations using structured tool calls³⁴. Enforce strict validation schemas on all outputs to guarantee that

LLM-generated payloads are parsed correctly before being passed to downstream API nodes⁵⁹.

Works cited

1. Understanding n8n from a System Design Perspective - DEV Community, <https://dev.to/tinhtinhcd/understanding-n8n-from-a-system-design-perspective-5075>
2. Build an AI Workflow Automation Platform Like n8n & Top 25 Alternatives - RBMSOFT, <https://rbmssoft.com/blogs/build-ai-workflow-automation-platform-like-n8n-and-top-25-alternatives/>
3. n8n Workflow Automation: Easy Guide for 2026 - Blck Alpaca, <https://blckalpaca.at/en/blog/n8n-workflow-automation-easy-guide-for-2026>
4. Low Code Architecture: A Comprehensive Guide - eSystems, <https://www.esystems.fi/en/blog/low-code-architecture-comprehensive-guide>
5. n8n vs Activepieces vs Windmill for Open Source Automation - Boolean & Beyond, <https://www.booleanbeyond.com/insights/n8n-vs-activepieces-vs-windmill-open-source-automation>
6. Comparison of n8n, ActivePieces, and Node-RED - Grokipedia, https://grokipedia.com/page/Comparison_of_n8n_ActivePieces_and_Node-RED
7. What is Windmill - Sixth Generation, <https://www.sixthgeneration.io/what-is-windmill>
8. Code in n8n | Build, <https://docs.n8n.io/build/code-in-n8n>
9. n8n - Yarn Classic, <https://classic.yarnpkg.com/en/package/n8n>
10. Self-Hosting | Definition and More - Activepieces Resources, <https://resources.activepieces.com/glossary/self-hosting>
11. How to Build No-Code AI Workflows Using Activepieces - freeCodeCamp, <https://www.freecodecamp.org/news/how-to-build-no-code-ai-workflows-using-activepieces/>
12. Changelog - Activepieces, <https://www.activepieces.com/docs/about/changelog>
13. Windmill | Code-first orchestration platform for internal software, <https://www.windmill.dev/>
14. Windmill | EuroStack Directory Project, <https://euro-stack.com/solutions/windmill>
15. Best Low-Code AI Workflow Automation Tools in 2026 - Firecrawl, <https://www.firecrawl.dev/blog/best-low-code-ai-workflow-automation-tools>
16. Top 10 Low-Code AI Workflow Automation Tools (2026) - Vellum, <https://www.vellum.ai/blog/top-low-code-ai-workflow-automation-tools>
17. Chocapikk/CVE-2026-21858: n8n Ni8mare - Unauthenticated Arbitrary File Read to RCE Chain (CVSS 10.0) - GitHub, <https://github.com/Chocapikk/CVE-2026-21858>
18. Enhancing ActivePieces Framework for Multi-Language Execution Via WebAssembly, <https://community.activepieces.com/t/enhancing-activepieces-framework-for-m>

- [ulti-language-execution-via-webassembly/6927](#)
19. Architecture and data exchange | Windmill,
<https://www.windmill.dev/docs/flows/architecture>
 20. Workflows as code - Windmill,
https://www.windmill.dev/docs/core_concepts/workflows_as_code
 21. What Is a Workflow Engine? - IBM,
<https://www.ibm.com/think/topics/workflow-engine>
 22. How n8n Works Internally - Architecture & Execution Engine Explained - C# Corner,
<https://www.c-sharpcorner.com/article/how-n8n-works-internally-architecture-execution-engine-explained/>
 23. Overview - Activepieces,
<https://www.activepieces.com/docs/install/architecture/overview>
 24. Visual Flow | React Flow Ready to Use templates | React Flow Examples | VisualFlow, <https://visualflow.dev/>
 25. Frontend Best Practices - Activepieces,
<https://www.activepieces.com/docs/handbook/engineering/playbooks/frontend-best-practices>
 26. Understand the database | Deploy - n8n Docs,
<https://docs.n8n.io/deploy/host-n8n/understand-the-architecture/understand-the-database>
 27. Environment Variables - Activepieces,
<https://www.activepieces.com/docs/install/reference/environment-variables>
 28. The AI Harness: why your AI coding agent is only as smart as the repo you put it in,
<https://dev.to/louaiboumediene/the-ai-harness-why-your-ai-coding-agent-is-only-as-smart-as-the-repo-you-put-it-in-cml>
 29. Data tables | Build - n8n Docs,
<https://docs.n8n.io/build/work-with-data/data-tables>
 30. Dynamic Data Table creation via JSON schema input with "Create Column If Not Exists",
<https://community.n8n.io/t/dynamic-data-table-creation-via-json-schema-input-with-create-column-if-not-exists/267689>
 31. Security | Deploy - n8n Docs,
<https://docs.n8n.io/deploy/host-n8n/configure-n8n/security>
 32. N8N Import Workflow JSON: Complete Guide + File Format Examples 2025 - Latenode Blog,
<https://latenode.com/blog/low-code-no-code-platforms/n8n-setup-workflows-self-hosting-templates/n8n-import-workflow-json-complete-guide-file-format-examples-2025>
 33. n8n Workflow JSON Schema Guide | PDF - Scribd,
<https://www.scribd.com/document/986524420/Reference>
 34. Inside n8n's AI Workflow Builder: A Complete Architecture Deep-Dive | by Rajveer Rathod,
<https://medium.com/@rajveer.rathod1301/inside-n8ns-ai-workflow-builder-a-co>

- [mplete-architecture-deep-dive-f2eeb2d57ec8](#)
35. Scaling n8n with Queue Mode: Automated Deployment Script | by The Minh Nguyen | NextGrowth.AI | Medium,
<https://medium.com/@nextgrowth.ai/scaling-n8n-with-queue-mode-automated-deployment-script-8ee5c317be6e>
 36. How to Scale n8n with Redis Queue Mode for Parallel Workflow Execution - Elestio blog,
<https://blog.elest.io/how-to-scale-n8n-with-redis-queue-mode-for-parallel-workflow-execution/>
 37. Trigger.dev vs BullMQ, <https://trigger.dev/vs/bullmq>
 38. Workers - Activepieces,
<https://www.activepieces.com/docs/install/architecture/workers>
 39. Autoscaling Queue Workers on Google Cloud Run (Redis/BullMQ) – How do you solve the HTTP-blindness? - n8n Community,
<https://community.n8n.io/t/autoscaling-queue-workers-on-google-cloud-run-redis-is-bullmq-how-do-you-solve-the-http-blindness/301961>
 40. CVE-2026-25049: Anatomy of an n8n Sandbox Escape, Forensics, and Enterprise Hardening - Penlilent,
<https://www.penlilent.ai/hackinglabs/cve-2026-25049-anatomy-of-an-n8n-sandbox-escape-forensics-and-enterprise-hardening/>
 41. n8n CVE Deep Dive: The Nervous System of Shadow AI and Its Critical Vulnerabilities,
<https://www.penlilent.ai/hackinglabs/n8n-cve-deep-dive-the-nervous-system-of-shadow-ai-and-its-critical-vulnerabilities/>
 42. VM module in Node.js - The Forgotten Power,
<https://2coffee.dev/en/articles/vm-module-in-nodejs-the-forgotten-power/>
 43. The Pitfalls of Node.js vm Module: Why It's Not a Sandbox - yceffort,
<https://yceffort.kr/en/2026/02/nodejs-deep-dive-sample>
 44. Critical Sandbox Escape in vm2 Enables Arbitrary Code Execution - Endor Labs,
<https://www.endorlabs.com/learn/cve-2026-22709-critical-sandbox-escape-in-v2-enables-arbitrary-code-execution>
 45. 2026: New N8N RCE Deep Dive into CVE-2026-25049 : r/netsec - Reddit,
https://www.reddit.com/r/netsec/comments/1qw2mw8/2026_new_n8n_rce_deep_dive_into_cve202625049/
 46. Eight for One: Multiple Vulnerabilities Fixed in the Node.js Runtime | Blog | Endor Labs,
<https://www.endorlabs.com/learn/eight-for-one-multiple-vulnerabilities-fixed-in-the-node-js-runtime>
 47. Why use Node.js vm module instead of isolated-vm for sandboxed LLM code execution? #17 - GitHub,
<https://github.com/universal-tool-calling-protocol/typescript-utcp/issues/17>
 48. We built Activepieces: open source alternative to Zapier (business automation) - Reddit,
https://www.reddit.com/r/selfhosted/comments/10pxkao/we_built_activepieces_open_source_alternative_to/

49. Enable SSRF protection | Deploy - n8n Docs,
<https://docs.n8n.io/deploy/host-n8n/configure-n8n/security/enable-ssrf-protection>
50. Server-Side Request Forgery (SSRF) in HTTP Node Leading to Internal Network Access,
<https://github.com/FlowiseAI/Flowise/security/advisories/GHSA-fvcw-9w9r-pxc7>
51. SSRF protection | Deploy - n8n Docs,
<https://docs.n8n.io/deploy/host-n8n/configure-n8n/basic-configuration/use-environment-variables/ssrf-protection>
52. Overview - React Flow, <https://reactflow.dev/examples>
53. AI Agent Architecture Patterns: Pick the Right Topology - n8n Blog,
<https://blog.n8n.io/ai-agent-architecture-patterns/>
54. Nested workflow with 'add table row' node working differently in v2 vs v1 - n8n Community,
<https://community.n8n.io/t/nested-workflow-with-add-table-row-node-working-differently-in-v2-vs-v1/248443>
55. How to build your first end-to-end AI workflow in n8n - The New Stack,
<https://thenewstack.io/build-n8n-ai-workflow/>
56. Build Custom AI Agents With Logic & Control | n8n Automation Platform,
<https://n8n.io/ai-agents/>
57. n8n AI Workflow Automation Guide: Build Smart Workflows in 2026 - DEV Community,
https://dev.to/nova_gg/n8n-ai-workflow-automation-guide-build-smart-workflows-in-2026-28cg
58. Generate SQL queries from schema only - AI-powered | n8n workflow template,
<https://n8n.io/workflows/2508-generate-sql-queries-from-schema-only-ai-powered/>
59. Structured Output Parser | Nodes - n8n Docs,
<https://docs.n8n.io/integrations/builtin/cluster-nodes/sub-nodes/n8n-nodes-langchain.outputparserstructured>